

# Natural Language Processing

---

SCHOOL OF INFOCOMM

# Need to know

- **Trainers**
  - Ho Chee Wai (ho\_chee\_wai@rp.edu.sg)
  - Tan Poh Keam (tan\_poh\_keam@rp.edu.sg)
- **Assessment**
  - Written assignment

# Learning Objectives

- Describe the principles, concepts and usage of NLP
- Perform text pre-processing activities
- Discuss and use suitable methods for feature engineering of text
- Discuss and use suitable approaches to analysis, classify and compare text
- Explain how deep-learning techniques are used for text generation tasks
- Consume Machine Learning services to build NLP applications
- Use Azure Bot Services to build conversation systems

# Where We're Headed

When	Topics
15 Jun	NLP Introduction Regular Expression Text Pre-processing
16 Jun	Features Engineering Document Classification (supervised learning)
17 Jun	Topic Modelling (unsupervised learning) Word Embedding
18 Jun	Topic Modelling (unsupervised learning) Text Generation
19 Jun	Google Language Services

When	Topics
22 Jun	Azure Bot Services Cognitive Services
23 Jun	Extra / extended classes (if needed)
25 Jun	Extra / extended classes (if needed)

# Introduction to NLP

---

# What is Natural Language Processing (NLP)

By “natural language” we mean a language that is used for everyday communication between humans.

Two main components:

- Natural Language Understanding (NLU)
- Natural Language Generation (NLG)

NLP is an Intersection of several fields

- Computer Science
- Artificial Intelligence
- Linguistics

NLP is AI-Complete

- Requires all types of knowledge & context awareness humans possess → It's hard

# Ambiguity of Language

- Synonymy – different words, same meaning
- Polysemy – same word, different meaning
- Text and speech are unstructured data
- No fixed structure – Sentence format
- No fixed schema – Grammar
- Misspell, slang, abbreviations

# History of NLP

- NLP has been through (at least) 3 major eras:
  - 1950s-1980s: Linguistics Methods and Rules
  - 1980s-Now: Statistical + Machine Learning Methods
  - Now - ??? : Deep Learning
- You're right near the start of a paradigm shift!



# 1950s - 1970s: Linguistics & Rules

- Approach focused on:
  - Linguistics: Grammar rules, sentence structure parsing
  - Handwritten Rules: Huge sets of logical (if/else) statements
  - Phase structure grammar : conversion of sentences into forms computers can understand
- Problems:
  - Too complex to maintain
  - Cannot scale
  - Cannot generalize

# 1980s - Now: Statistical / Machine Learning Methods

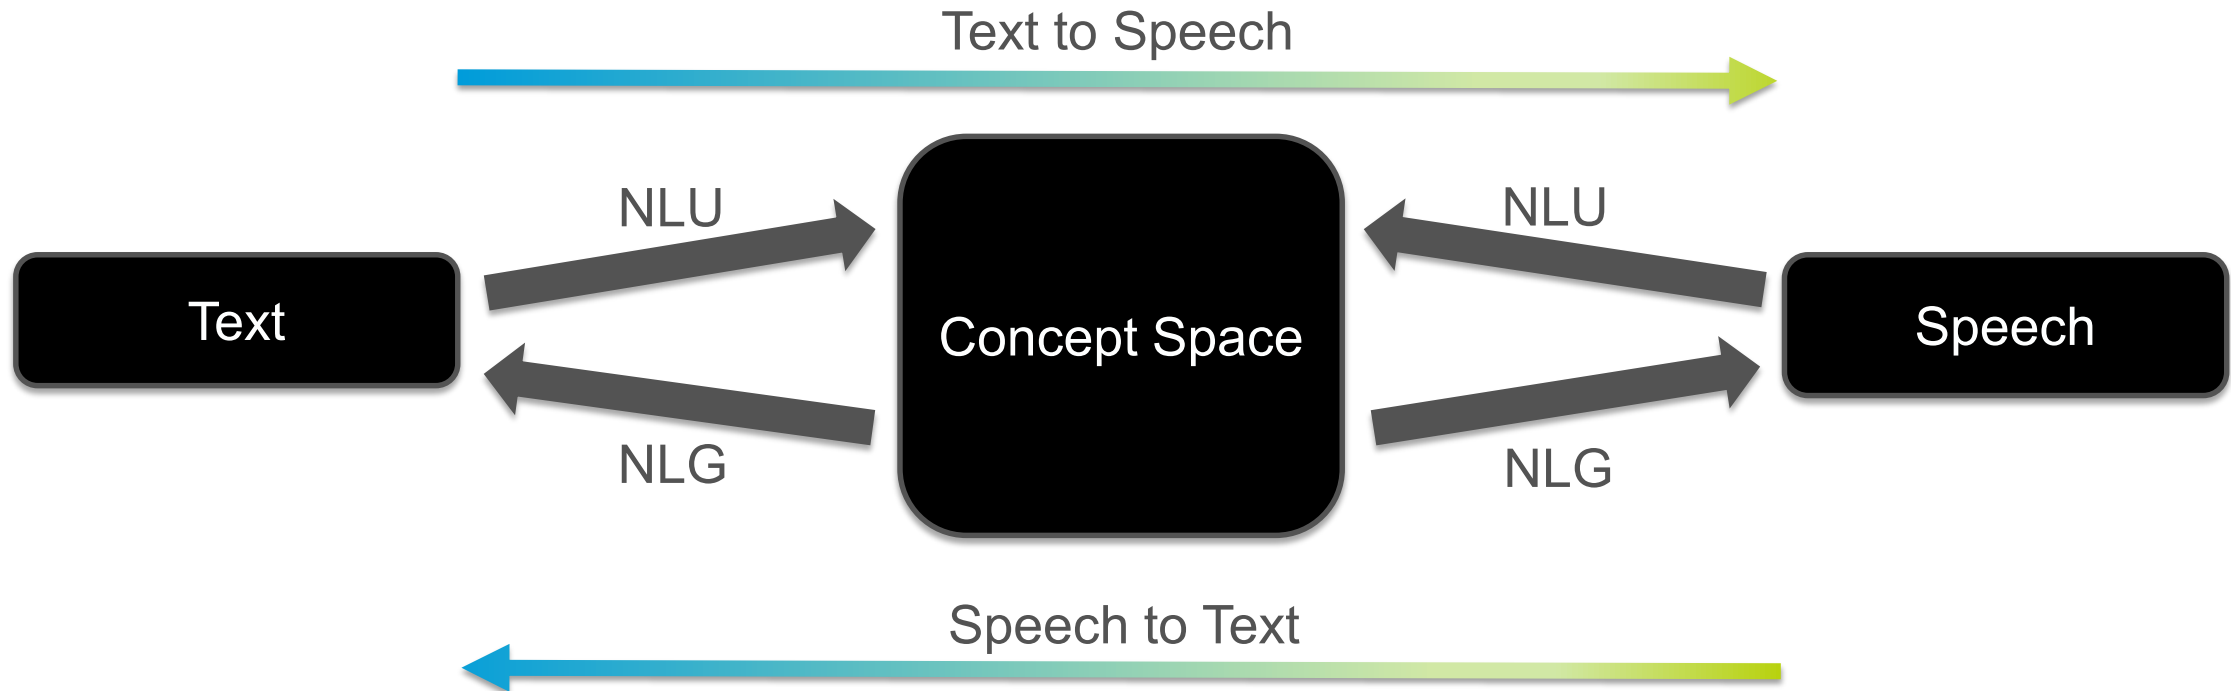
- Approach shifted from linguistics to data-driven
- Increasing computational power and ease to access of text
  - First web page (1991) → discussion forums, blogs, news portal
  - Digital archives
- NLP starts using statistical and probabilistic models
  - Data mining → text mining
- Generic machine learning algorithms applied to NLP tasks
  - Sentiment analysis using logistic regression
  - Language models with Markov models

# Now and future: Deep Learning

- More advances in computing power with parallelization (GPU)
- Availability of large datasets becomes the norm
- Neural Network
  - Learnt word representation with finite dimensions
  - Capture semantic and relationships among words
- RNN / LSTM
  - Allows sequential processing and learning of text
  - Application into machine translation tasks and questions/answering systems
- Attention-based model
  - A way to place various degree of focus (attention) on different part of the text
  - Break-through in machine translation and text generation tasks

# NLP: Speech vs Text

- Data source can refer to written text or speech
- Goal of both is the same: translate raw data (text or speech) into underlying concepts (NLU) then possibly into the other form (NLG)



# NLU Applications

- Classification
- Natural Language Search
- Document Recommendation
- Topic Modeling
- Language Identification
- Intent matching

# NLU Application: Document Classification

- Classify documents (discrete collections of text) - into categories
  - Classify emails as spam vs not spam
  - Classify product reviews as positive vs negative
  - Assign labels to documents

# NLU Application: Document Recommendation

- Choosing the most relevant document based on some information or "finger print":
  - Show most relevant webpages based on query to search engine
  - Recommend news articles based on past articles liked or read
  - Recommend restaurants based on restaurant reviews

# NLU Application: Topic Modeling

Breaking a set of documents into topics at the word level

- See how prevalence of certain topics covered in a magazine changes over time
- Find documents belonging to a certain topic

## Election 2015: Parties row over GP out-of-hours cover

3 hours ago | Election 2015



Figures suggest almost 600 fewer GP surgeries in England open at evenings and weekends than before 2010, Labour has claimed.

Health spokesman Andy Burnham said the coalition had created queues outside practices and diverted people to A&E.

Tory Health Secretary Jeremy Hunt said Labour's numbers were wrong and that out-of-hours cover was being extended.

The Lib Dems also said Labour's figures - obtained through a **parliamentary question** - were out of date.

Mr Burnham announced the analysis as his party unveiled a new poster, which reworks the Conservatives' "Labour isn't working" image of 1979 by depicting a huge queue outside a waiting room with the title: "The doctor can't see you now."

Politics or Health?

## Chelsea FC reports a record £18m in annual profit

13 November 2014 | Business



Chelsea Football Club has reported a record profit of £18.4m (\$29m) for the year to June 2014 - despite last season's lack of silverware.

The London team has only once before turned a profit in the 10 years since it was acquired by Russian billionaire Roman Abramovich.

**The club said a new TV broadcasting deal**, as well as the sale of players such as Juan Mata, had boosted profits.

Chelsea are currently unbeaten at the top of the Premier League.

### Business of Sport

**Premier League to share £1bn**

**Adidas unveils Europe factory plan**

**Fans to protest over ticket prices**

**Controversial Gatlin gets Nike deal**

Business or Sport?



# NLU Application: Intent Matching

Understanding that there are many ways to say or ask for the same thing

- Use in dialog systems

*May I know the opening time for the store?*

*Can I come to shop now?*

*What time do you close?*

*Is the mall still open?*

*I can come at 8 pm, is it okay ?*

# NLG Applications

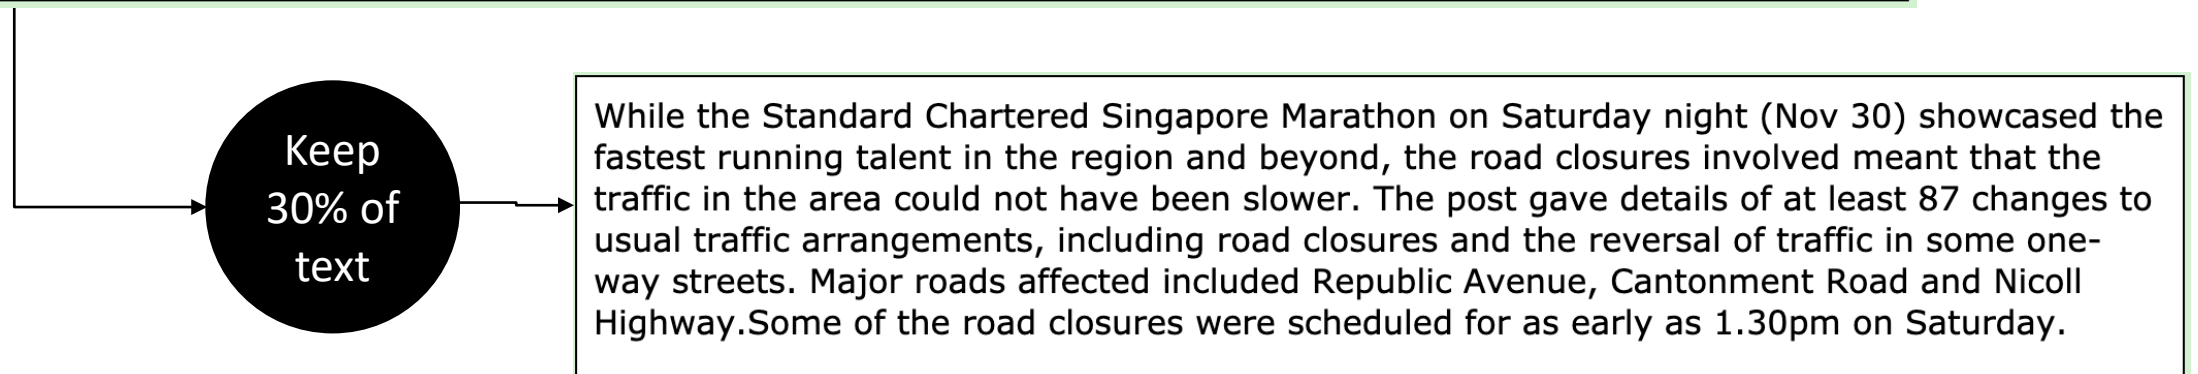
- Machine Translation
- Document Summarization
- Text Generation
- Question Answering
- Image Captioning

Notice NLU is almost a prerequisite for NLG

# NLG Application: Document Summarization

## Read less

While the Standard Chartered Singapore Marathon on Saturday night (Nov 30) showcased the fastest running talent in the region and beyond, the road closures involved meant that the traffic in the area could not have been slower. Unlike the past 17 times when the annual event was flagged off in the morning and roads were less busy, Saturday's marathon started at 6pm as organisers sought to raise the event's profile. The changed timing and heavy traffic conditions around that hour led to traffic gridlock and left some drivers furious. While thousands of runners pounded the asphalt, motorists interviewed by TODAY were stuck in traffic for up to 2.5 hours. For some couples holding wedding banquets in town on Saturday, as many as half their guests arrived late, with many no-shows who turned back in frustration. On Nov 27, the police took to Facebook and Twitter to warn commuters of delays during the event. The post gave details of at least 87 changes to usual traffic arrangements, including road closures and the reversal of traffic in some one-way streets. Major roads affected included Republic Avenue, Cantonment Road and Nicoll Highway. Some of the road closures were scheduled for as early as 1.30pm on Saturday. Motorists who had to travel around the affected areas for work, such as wedding planners and performers, were also caught up in the bumper-to-bumper traffic.



# NLG Application: Document Summarization

Automatically generate text summaries of documents

<b>Input: Article 1st sentence</b>	<b>Model-written headline</b>
metro-goldwyn-mayer reported a third-quarter net loss of dlr 16 million due mainly to the effect of accounting rules adopted this year	mgm reports 16 million net loss on higher revenue
starting from july 1, the island province of hainan in southern china will implement strict market access control on all incoming livestock and animal products to prevent the possible spread of epidemic diseases	hainan to curb spread of diseases
australian wine exports hit a record 52.1 million liters worth 260 million dollars (143 million us) in september, the government statistics office reported on monday	australian wine exports hit record high in september

Source: <https://ai.googleblog.com/2016/08/text-summarization-with-tensorflow.html>

# NLG Application: Machine Translation

Automatically translate text between languages

The diagram illustrates a machine translation process. On the left, a text box shows the input English sentence: "How can I get the hospital. My son just slipped and have a cut on his chin. He is bleeding badly". Above this text are three tabs: "ENGLISH - DETECTED" (selected), "ENGLISH", and "SPANIS". A large green arrow points from this box to the right. On the right, there are two stacked translation boxes. The top box shows the Chinese (Simplified) translation. It has three tabs at the top: "JAPANESE", "CHINESE (SIMPLIFIED)" (selected), and "ENG". The Chinese text is "我怎么去医院。我儿子刚刚滑倒，下巴上有伤口。他流血严重". Below it is the Pinyin: "Wǒ zěnmē qù yīyuàn. Wǒ érzi gānggāng huá dǎo, xiàbā shàng yǒu shāngkǒu. Tā liúxuè yánzhòng". The bottom box shows the Malay translation. It has three tabs at the top: "MALAY" (selected), "JAPANESE", and "CHINESE (SIMPLIFIED)". The Malay text is "Bagaimana saya boleh mendapatkan hospital. Anak saya hanya tergelincir dan memotong dagunya. Dia berdarah dengan".

ENGLISH - DETECTED   ENGLISH   SPANIS

How can I get the hospital. My son just slipped and have a cut on his chin. He is bleeding badly

JAPANESE   CHINESE (SIMPLIFIED)   ENG

我怎么去医院。我儿子刚刚滑倒，下巴上有伤口。他流血严重

Wǒ zěnmē qù yīyuàn. Wǒ érzi gānggāng huá dǎo, xiàbā shàng yǒu shāngkǒu. Tā liúxuè yánzhòng

MALAY   JAPANESE   CHINESE (SIMPLIFIED)

Bagaimana saya boleh mendapatkan hospital. Anak saya hanya tergelincir dan memotong dagunya. Dia berdarah dengan

Example from Google Translate

Reading : <https://ai.googleblog.com/2016/09/a-neural-network-for-machine.html>

# NLG Application: Text Generation

SYSTEM PROMPT  
(HUMAN-WRITTEN)

*In a shocking finding, scientist discovered a herd of unicorns living in a remote, previously unexplored valley, in the Andes Mountains. Even more surprising to the researchers was the fact that the unicorns spoke perfect English.*

MODEL  
COMPLETION  
(MACHINE-  
WRITTEN, 10 TRIES)

The scientist named the population, after their distinctive horn, Ovid's Unicorn. These four-horned, silver-white unicorns were previously unknown to science.

Now, after almost two centuries, the mystery of what sparked this odd phenomenon is finally solved.

Dr. Jorge Pérez, an evolutionary biologist from the University of La Paz, and several companions, were exploring the Andes Mountains when they found a small valley, with no other animals or humans. Pérez noticed that the valley had what appeared to be a natural fountain, surrounded by two peaks of rock and silver snow.

# NLG Application: Question and Answering

QnA is concerned with building systems that automatically **answer questions** posed by humans in a natural language.

Question	Generated Answer	Correct	Probability
Who wrote the book the origin of species?	Charles Darwin	✓	83.4%
Who is the founder of the ubuntu project?	Mark Shuttleworth	✓	82.0%
Who is the quarterback for the green bay packers?	Aaron Rodgers	✓	81.1%
Panda is a national animal of which country?	China	✓	76.8%
Who came up with the theory of relativity?	Albert Einstein	✓	76.4%
When was the first star wars film released?	1977	✓	71.4%
What is the most common blood type in sweden?	A	✗	70.6%
Who is regarded as the founder of psychoanalysis?	Sigmund Freud	✓	69.3%
Who took the first steps on the moon in 1969?	Neil Armstrong	✓	66.8%
Who is the largest supermarket chain in the uk?	Tesco	✓	65.3%
What is the meaning of shalom in english?	peace	✓	64.0%
Who was the author of the art of war?	Sun Tzu	✓	59.6%
Largest state in the us by land mass?	California	✗	59.2%
Green algae is an example of which type of reproduction?	asexual reproduction	✓	56.5%



# NLG

## Application: Question Answering

Source:  
<https://venturebeat.com/2019/11/25/amazon-researchers-use-nlp-data-set-to-improve-alexa-answers/>



Image Credit: Amazon

Improving the quality of voice assistants' responses to questions is of interest to tech giants like Google, Apple, and Microsoft, who seek to address shortfalls in their respective natural language processing (NLP) technologies. They've plenty in the way of motivation — more than 50% of U.S. smart speaker owners say they ask questions of their devices, according to a survey conducted last year by Adobe.



# NLG Application: Image Captioning

azure.microsoft.com/en-us/services/cognitive-services/computer-vision/

Product overview

Features

Security

Pricing

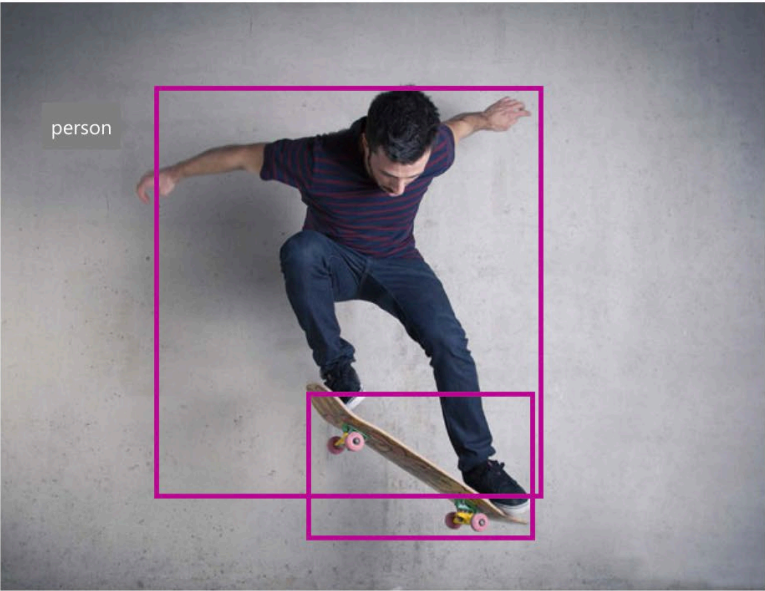
Documentation

Customer stories

FAQs

Free account >


person



FEATURE NAME:	VALUE
Objects	[ { "rectangle": { "x": 238, "y": 299, "w": 177, "h": 117 }, "object": "Skateboard", "confidence": 0.903 }, { "rectangle": { "x": 118, "y": 63, "w": 305, "h": 321 }, "object": "person", "confidence": 0.955 } ]
Tags	[ { "name": "skating", "confidence": 0.999951541 }, { "name": "snowboarding", "confidence": 0.990067363 }, { "name": "sports equipment", "confidence": 0.9774853 }, { "name": "person", "confidence": 0.9605776 }, { "name": "roller skating", "confidence": 0.945730746 }, { "name": "boardsport", "confidence": 0.9242261 }, { "name": "man",

Image URL

Submit


 Browse

Demo:

- <https://azure.microsoft.com/en-us/services/cognitive-services/computer-vision/>

25

# Publications on Applied Areas of NLP



Credit: A Survey of the Usages of Deep Learning for Natural Language Processing , Daniel W. Otter, Julian R. Medina, and Jugal K. Kalita (<https://arxiv.org/pdf/1807.10854.pdf>)

# NLP Benchmarks

## **BLEU Score - BiLingual Evaluation Understudy**

- Metric to evaluate translations from one language to another

## **SQuAD — Stanford Question Answering Dataset**

- Dataset of text and related questions with answers from the dataset. Given a question and the model needs to predict the right answer

## **MS MACRO — MACHine Reading COMprehension Dataset**

- Large scale dataset focused on machine reading comprehension

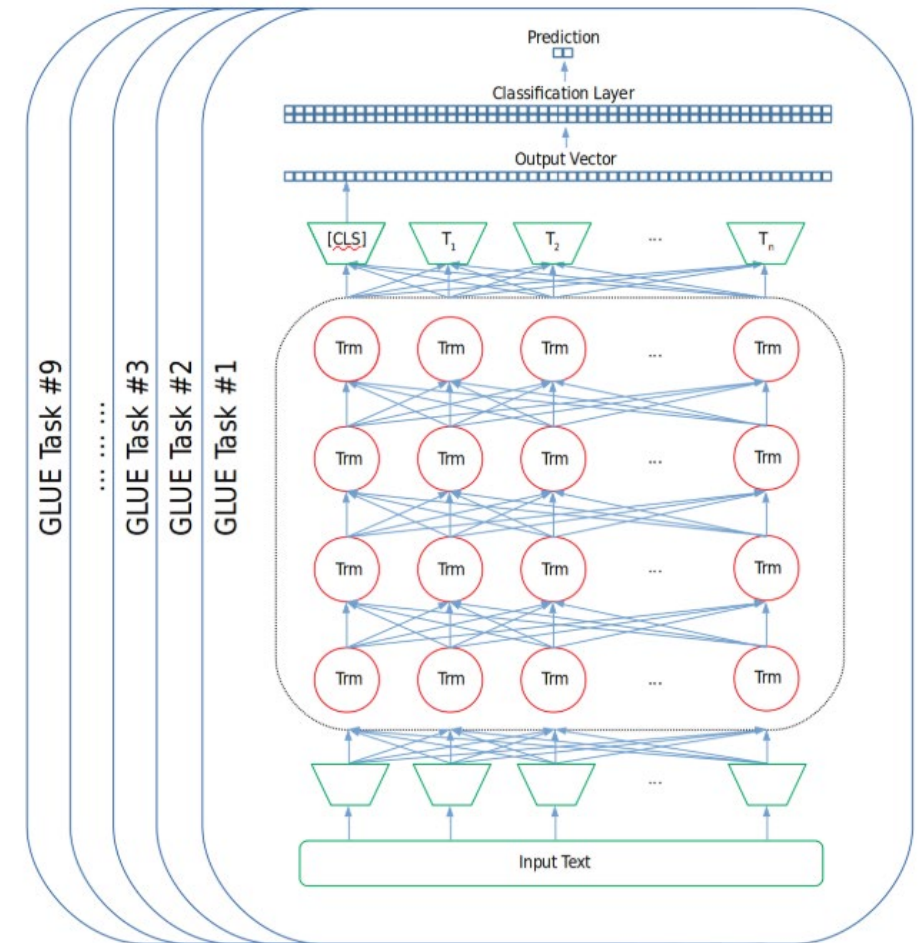
## **GLUE - General Language Understanding Evaluation**

- Evaluates the performance of a model on a collection of multiple NLP tasks

# GLUE

- A collection of resources for training, evaluating, and analysing natural language understanding systems
- **Language understanding tasks** built on established **existing datasets** and selected to cover a diverse range of **dataset sizes**, **text genres**, and **degrees of difficulty**
- Website contains a leader board for **tracking performance of emerging models / algorithms** against benchmark (<https://gluebenchmark.com/leaderboard>)

$$\sum \text{Individual Task Scores} = \text{Final GLUE Score}$$



Credit: <https://mccormickml.com/2019/11/05/GLUE/>

Dataset	Description	Data example
CoLA	Is the sentence grammatical or ungrammatical?	"This building is than that one." = <b>Ungrammatical</b>
SST-2	Is the movie review positive, negative, or neutral?	"The movie is funny , smart , visually inventive , and most of all , alive ." = <b>.93056 (Very Positive)</b>
MRPC	Is the sentence B a paraphrase of sentence A?	A) "Yesterday , Taiwan reported 35 new infections , bringing the total number of cases to 418 ." B) "The island reported another 35 probable cases yesterday , taking its total to 418 ." = <b>A Paraphrase</b>
STS-B	How similar are sentences A and B?	A) "Elephants are walking down a trail." B) "A herd of elephants are walking along a trail." = <b>4.6 (Very Similar)</b>
QQP	Are the two questions similar?	A) "How can I increase the speed of my internet connection while using a VPN?" B) "How can Internet speed be increased by hacking through DNS?" = <b>Not Similar</b>
MNLI-mm	Does sentence A entail or contradict sentence B?	A) "Tourist Information offices can be very helpful." B) "Tourist Information offices are never of any help." = <b>Contradiction</b>
QNLI	Does sentence B contain the answer to the question in sentence A?	A) "What is essential for the mating of the elements that create radio waves?" B) "Antennas are required by any radio receiver or transmitter to couple its electrical connection to the electromagnetic field." = <b>Answerable</b>
RTE	Does sentence A entail sentence B?	A) "In 2003, Yunus brought the microcredit revolution to the streets of Bangladesh to support more than 50,000 beggars, whom the Grameen Bank respectfully calls Struggling Members." B) "Yunus supported more than 50,000 Struggling Members." = <b>Entailed</b>
WNLI	Sentence B replaces sentence A's ambiguous pronoun with one of the nouns - is this the correct noun?	A) "Lily spoke to Donna, breaking her concentration." B) "Lily spoke to Donna, breaking Lily's concentration." = <b>Incorrect Referent</b>



# Cognitive Language Services for Building Apps

[Overview](#)[Features](#)[APIs](#)[Customer stories](#)[Documentation](#)[FAQs](#)[Decision](#)[Language](#)[Speech](#)[Vision](#)[Web search](#)

Extract meaning from unstructured text

**Immersive Reader** PREVIEW

Help readers of all abilities comprehend text using audio and visual cues.

**Language Understanding**

Build natural language understanding into apps, bots, and IoT devices.

**QnA Maker**

Create a conversational question and answer layer over your data.

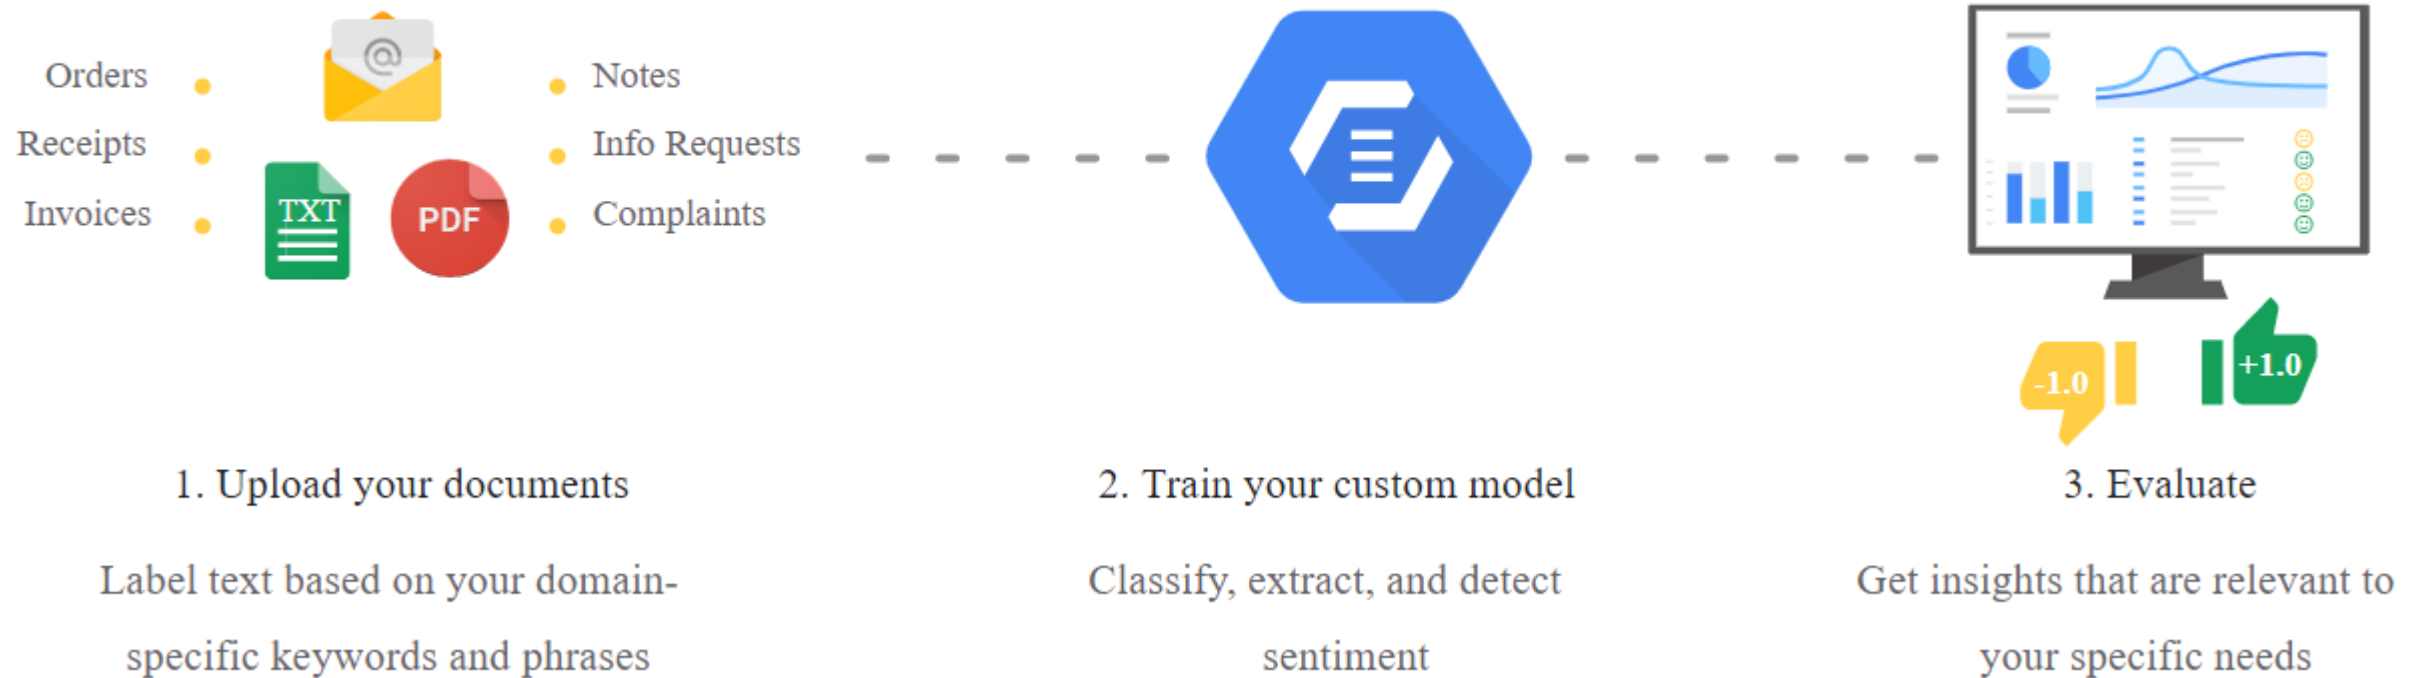
**Text Analytics**

Detect sentiment, key phrases, and named entities.

**Translator**

Detect and translate more than 60 supported languages.

# Language Services for Building Apps



# Key Takeaways

- It is hard to understand language, but tech and research community are discovering better way to create better algorithms and models to perform NLP tasks
- There are many NLP enabled applications around us
- You can create models or consume language API or services